


PoreAccess-CO₂: From Random Packing to an Interpretable Transport Closure

Structure-resolved diffusion and adsorption in
DEM-generated zeolite beds

Technical University of Munich
Chair of Process Systems Engineering

Ali Khalifa & Heiko Briesen



2026-09-12

Introduce the project: how random packing geometry can inform a reduced model of transient adsorption.



Presentation outline

Motivation and research question

Pore-network reference model

Reduction and diffusivity prediction

Results and robustness

Conclusions and next steps

We first motivate inlet accessibility, then explain the PNM reference and its mean uptake output. Next we reduce it to a homogeneous bed and develop the diffusivity closure. Results and conclusions follow.

Section 1 of 5

Motivation and research question

Why should inlet-connected pathways matter?

Why should inlet-connected pathways matter?.

Why examine access from the inlet?

What porosity misses

- Porosity measures how much void space exists.
- It does not describe bottlenecks, path lengths or parallel connections.
- Similar void fractions can therefore hide different gas-supply pathways.

Application: CO₂ adsorption in zeolite 13X, used here as a controlled model problem.

Hypothesis: inlet-connected diffusive pathways contain information that porosity misses.

Why the inlet is the reference

- Gas enters from a reservoir below the bed.
- Adsorption removes gas as it diffuses into the packing.
- Access from the inlet may therefore help explain transient uptake differences.

2026-09-12

└ Motivation and research question

└ Why examine access from the inlet?

The motivation is transport through random geometry. Two packings can have similar porosity but very different arrangements of constrictions and parallel pathways. In our experiment gas is supplied at the bottom and removed by adsorption as it enters. That makes the inlet a physically meaningful reference boundary. This leads us to test inlet-to-interior accessibility as a descriptor. It is a hypothesis, not a claim that the inlet alone controls every time scale. CO₂ and zeolite provide a controlled application for studying this question.

Why examine access from the inlet?

What porosity misses

- Porosity measures how much void space exists.
- It does not describe bottlenecks, path lengths or parallel connections.
- Similar void fractions can therefore hide different gas-supply pathways.

Application: CO₂ adsorption in zeolite 13X, used here as a controlled model problem.

Hypothesis: inlet-connected diffusive pathways contain information that porosity misses.


Why the inlet is the reference

- Gas enters from a reservoir below the bed.
- Adsorption removes gas as it diffuses into the packing.
- Access from the inlet may therefore help explain transient uptake differences.


Same porosity can hide different inlet-connected pathways

Can a pore-accessibility measure predict bed-scale transient transport better than porosity alone?

Packing A: inlet bottleneck



Packing B: multiple inlet paths




- pore body
- diffusive throat
- zeolite bead
- inlet bottleneck

Same beads and similar porosity; different numbers and conductances of inlet-connected paths.

2026-09-12

└ Motivation and research question

└ Same porosity can hide different inlet-connected pathways



This schematic introduces the central problem. The two beds contain the same monodisperse beads and can have almost the same total porosity. Their inlet-connected paths can nevertheless differ strongly. Packing A illustrates a network reached through one restrictive inlet bottleneck. Packing B illustrates several paths that connect the reservoir to the bed interior.

The blue lines show only simplified inlet-connected paths, not every throat in a real extracted network. The sketches are therefore hypotheses rather than simulation results. They motivate our research question: can a quantitative pore-accessibility measure predict transient bed-scale transport better than bulk porosity alone?

Section 2 of 5

Pore-network reference model

From particle geometry to the mean uptake curve

From particle geometry to the mean uptake curve.

Twenty beds isolate variability caused by random arrangement



Quantity	Fixed value
beads per bed	1850
bead diameter	1.6 mm
tube diameter	16 mm
D_{tube}/d_p	10
temperature	298.15 K
inlet c_{CO_2}	6.05 mol m ⁻³
simulation time	5000 s

Controlled diffusion–adsorption problem

- same beads and tube; only the insertion seed changes
- dry, isothermal gas supplied from below
- diffusion and adsorption, with no imposed flow
- heat effects and competitive water adsorption excluded

$c_{\text{in}} = y_{\text{CO}_2} p / (RT) \approx 6.05 \text{ mol m}^{-3}$ at 15 vol.%, 1 bar and 298 K.
15% represents coal-combustion flue gas (typical CO₂: 10–15%).

A dry, cooled reference gas fixes the driving concentration across all beds.

Source: U.S. DOE, Quadrennial Technology Review (2015), Sec. 4.E.

The ensemble tests arrangement effects under otherwise identical conditions.

2026-09-12

- └ Pore-network reference model
- └ Twenty beds isolate variability caused by random arrangement

Twenty beds isolate variability caused by random arrangement	
Quantity	Fixed value
beads per bed	1850
bead diameter	1.6 mm
tube diameter	16 mm
D_{tube}/d_p	10
temperature	298.15 K
inlet c_{CO_2}	6.05 mol m ⁻³
simulation time	5000 s

Controlled diffusion–adsorption problem

- same beads and tube; only the insertion seed changes
- dry, isothermal gas supplied from below
- diffusion and adsorption, with no imposed flow
- heat effects and competitive water adsorption excluded

$c_{\text{in}} = y_{\text{CO}_2} p / (RT) \approx 6.05 \text{ mol m}^{-3}$ at 15 vol.%, 1 bar and 298 K.
15% represents coal-combustion flue gas (typical CO₂: 10–15%).
A dry, cooled reference gas fixes the driving concentration across all beds.
Source: U.S. DOE, Quadrennial Technology Review (2015), Sec. 4.E.

The ensemble tests arrangement effects under otherwise identical conditions.

The inlet value corresponds to a CO2 mole fraction of approximately 15 percent at one bar: c equals y times p divided by RT . The DOE report gives a typical coal-combustion flue-gas range of 10 to 15 percent CO2, which motivates this representative concentration. Temperature is held at 298.15 K and the gas is dry, so this is a cooled, simplified reference condition. We do not simulate hot, wet industrial flue gas or its other constituents. Keeping the same reservoir concentration for all twenty beds isolates arrangement effects. The cited concentration range motivates the boundary condition, not a claim of industrial process validation.

Pore-network geometry and throat diffusion

Network construction

- four bead centres define a Delaunay tetrahedron; its valid circumcentre represents a pore node
- tetrahedra sharing a triangular face are joined by a throat
- topology comes from particle coordinates, not voxels
- 2^{22} Sobol points sample the solid and gas space.
- Gas-point counts allocate the calculated void volume to pores.

Diffusion through throat $i-j$

$$G_{ij} = D_m \frac{A_{ij}}{L_{ij}}, \quad \dot{n}_{ij} = G_{ij}(c_j - c_i).$$

$D_m = 1.5 \times 10^{-5} \text{ m}^2 \text{ s}^{-1}$: molecular CO_2 diffusivity in the gas.

$L_{ij} = \|\mathbf{x}_j - \mathbf{x}_i\|$ is the pore-centre distance. The irregular opening between three beads is approximated by an equivalent circular area

$$A_{ij} = \pi r_{ij}^2,$$

where r_{ij} is the clearance of the shared Delaunay face.

The circular throat area is a geometric modelling assumption; G_{ij} has units of $\text{m}^3 \text{ s}^{-1}$.

2026-09-12

- └ Pore-network reference model
- └ Pore-network geometry and throat diffusion

Pore-network geometry and throat diffusion

Network construction

- four bead centres define a Delaunay tetrahedron; its valid circumcentre represents a pore node
- tetrahedra sharing a triangular face are joined by a throat
- topology comes from particle coordinates, not voxels
- 2^{22} Sobol points sample the solid and gas space.
- Gas-point counts allocate the calculated void volume to pores.

Diffusion through throat $i-j$

$G_{ij} = D_m \frac{A_{ij}}{L_{ij}}, \quad \dot{n}_{ij} = G_{ij}(c_j - c_i).$

$D_m = 1.5 \times 10^{-5} \text{ m}^2 \text{ s}^{-1}$: molecular CO_2 diffusivity in the gas.

$L_{ij} = \|\mathbf{x}_j - \mathbf{x}_i\|$ is the pore-centre distance. The irregular opening between three beads is approximated by an equivalent circular area

$A_{ij} = \pi r_{ij}^2,$

where r_{ij} is the clearance of the shared Delaunay face.

The circular throat area is a geometric modelling assumption; G_{ij} has units of $\text{m}^3 \text{ s}^{-1}$.

The particle centres define Delaunay tetrahedra. Valid circumcentres become pore nodes, and a shared triangular face defines a candidate throat. $D_{\text{sub m}}$ is the prescribed molecular diffusivity of CO_2 in the gas, 1.5 times ten to the minus five square metres per second. It is a fixed model input, not the fitted bed-scale effective diffusivity. Conductance equals this diffusivity times the approximate opening area divided by pore-centre distance. Sobol sampling assigns gas volumes to the nodes. Detailed construction sketches remain in the optional appendix.

The pore network solves diffusion and adsorption together

Gas balance for each pore i

$$V_i \frac{dc_i}{dt} = \sum_{j \in \mathcal{N}(i)} G_{ij}(c_j - c_i) + G_{iR}(c_R - c_i) - S_i.$$

- V_i : gas volume, c_i : gas concentration
- G_{ij} : throat conductance [$\text{m}^3 \text{s}^{-1}$]
- G_{iR} : connection to reservoir at c_R , zero away from the inlet
- S_i : net gas transfer to adjacent beads [mol s^{-1}]

Uptake for each bead k

$$\frac{dq_k}{dt} = k_{\text{LDF}}(q_k^* - q_k).$$

- q_k : current loading [mol kg^{-1}]
- k_{LDF} [s^{-1}]: linear-driving-force rate coefficient
- $1/k_{\text{LDF}}$: relaxation time at fixed gas concentration

One coupled system: one concentration per pore and one loading per bead.

2026-09-12

└ Pore-network reference model

└ The pore network solves diffusion and adsorption together

The pore network solves diffusion and adsorption together

Gas balance for each pore i

$$V_i \frac{dc_i}{dt} = \sum_{j \in \mathcal{N}(i)} G_{ij}(c_j - c_i) + G_{iR}(c_R - c_i) - S_i.$$

- V_i : gas volume, c_i : gas concentration
- G_{ij} : throat conductance [$\text{m}^3 \text{s}^{-1}$]
- G_{iR} : connection to reservoir at c_R , zero away from the inlet
- S_i : net gas transfer to adjacent beads [mol s^{-1}]

Uptake for each bead k

$$\frac{dq_k}{dt} = k_{\text{LDF}}(q_k^* - q_k).$$

- q_k : current loading [mol kg^{-1}]
- k_{LDF} [s^{-1}]: linear-driving-force rate coefficient
- $1/k_{\text{LDF}}$: relaxation time at fixed gas concentration

One coupled system: one concentration per pore and one loading per bead.

This is the transient diffusion equation on the network, expressed as a balance on each pore volume. The first term exchanges gas with neighbouring pores, the second connects inlet pores to the finite-conductance reservoir, and the final term transfers gas to beads. Each bead has one loading equation even if several pores surround it. LDF means linear driving force: the rate is proportional to the difference between equilibrium and current loading. All pore concentrations and bead loadings are integrated together from the prescribed initial state. The total transfer satisfies $\sum_i S_i = \sum_k m_k dq_k/dt$, so adsorption transfers gas into the solid inventory. In the code, positive uptake is allocated among incident pores using incidence weights multiplied by positive gas concentration, preventing empty pores from supplying gas. Desorption uses normalized incidence weights. No independent gas-flow or pressure equation is solved.

Local equilibrium sets the adsorption target

Local equilibrium loading

$$p_k = c_k RT, \quad q_k^* = \frac{q_s b p_k}{1 + b p_k}.$$

- c_k : average gas concentration around bead k
- p_k : local CO₂ partial pressure [Pa]
- q_k^* : equilibrium target [mol kg⁻¹]

Fixed material parameters

- q_s : maximum CO₂ loading [mol kg⁻¹]
- b : Langmuir affinity [Pa⁻¹], so $b p_k$ is dimensionless
- k_{LDF} : kinetic coefficient [s⁻¹]

Actual loading evolves in time

$$\frac{dq_k}{dt} = k_{\text{LDF}}(q_k^* - q_k).$$

Local gas concentrations set the target. Kinetics determine how quickly beads approach it.

2026-09-12

└ Pore-network reference model

└ Local equilibrium sets the adsorption target

Local equilibrium sets the adsorption target

Local equilibrium loading

$$p_k = c_k RT, \quad q_k^* = \frac{q_s b p_k}{1 + b p_k}.$$

Local gas concentrations set the target. Kinetics determine how quickly beads approach it.

Fixed material parameters

- q_s : maximum CO₂ loading [mol kg⁻¹]
- b : Langmuir affinity [Pa⁻¹], so $b p_k$ is dimensionless
- k_{LDF} : kinetic coefficient [s⁻¹]

Actual loading evolves in time

$$\frac{dq_k}{dt} = k_{\text{LDF}}(q_k^* - q_k).$$

The local gas concentration around bead k is a normalized, pore-volume-weighted average over its associated pores. The implementation uses nonnegative concentrations in this calculation. The ideal-gas relation gives the local CO₂ partial pressure and the Langmuir isotherm supplies the equilibrium loading. This equilibrium target is not a new differential unknown. The actual loading evolves through the LDF equation introduced on the preceding slide. The saturation loading and affinity are prescribed material parameters, and k_{LDF} is fixed across packings. Adsorption removes gas from the incident pores and adds it to the bead, preserving the combined inventory apart from reservoir exchange.

The pore network delivers the reference uptake curve

Resolved transient outputs

- Gas concentration $c_i(t)$ in every pore
- Adsorbed loading $q_k(t)$ of every bead
- Different bead loadings reveal incomplete utilization

The detailed model retains the actual pore connections and bead associations.

Bed-average uptake

$$\bar{q}_{\text{PNM}}(t) = \frac{\sum_k m_k q_k(t)}{\sum_k m_k}.$$

For identical beads:

$$\bar{q}_{\text{PNM}}(t) = \frac{1}{N_b} \sum_{k=1}^{N_b} q_k(t).$$

This curve is the reference for fitting the reduced 1D model.

Next: can a homogeneous 1D bed reproduce this complete uptake curve?

2026-09-12

└ Pore-network reference model

└ The pore network delivers the reference uptake curve

The pore network delivers the reference uptake curve

Resolved transient outputs

- Gas concentration $c_i(t)$ in every pore
- Adsorbed loading $q_k(t)$ of every bead
- Different bead loadings reveal incomplete utilization

The detailed model retains the actual pore connections and bead associations.

Bed-average uptake

$$\bar{q}_{\text{PNM}}(t) = \frac{\sum_k m_k q_k(t)}{\sum_k m_k}$$

For identical beads:

$$\bar{q}_{\text{PNM}}(t) = \frac{1}{N_b} \sum_{k=1}^{N_b} q_k(t).$$

This curve is the reference for fitting the reduced 1D model.

Next: can a homogeneous 1D bed reproduce this complete uptake curve?

The network calculation produces many concentration and loading histories. For model reduction, we summarize the adsorbed inventory as the mass-weighted mean loading of all beads. Because the beads have equal masses, this is their arithmetic mean. We retain the full time series, not only the final loading. The PNM curve is a numerical reference under our modelling assumptions, not an experimental measurement. We now ask whether a much smaller homogeneous model can reproduce this curve.

Section 3 of 5

Reduction and diffusivity prediction

Fit a reduced model, then predict its coefficient

Fit a reduced model, then predict its coefficient.

Reduce the detailed network to a homogeneous 1D bed

Detailed pore-network model

- Actual pores and throats from each packing
- One gas concentration per pore
- One loading per bead
- Reference mean uptake $\bar{q}_{\text{PNM}}(t)$

Homogeneous 1D model

- Replace the network by 100 equal axial slices
- Each slice spans the full tube cross-section
- One gas concentration and one mean solid loading per slice
- Represent unresolved pathways with a single diffusivity D

Fit D so the reduced model reproduces the network's mean uptake history.

2026-09-12

└ Reduction and diffusivity prediction

└ Reduce the detailed network to a homogeneous 1D bed

Reduce the detailed network to a homogeneous 1D bed

Detailed pore-network model

- Actual pores and throats from each packing
- One gas concentration per pore
- One loading per bead
- Reference mean uptake $\bar{q}_{\text{PNM}}(t)$

Homogeneous 1D model

- Replace the network by 100 equal axial slices
- Each slice spans the full tube cross-section
- One gas concentration and one mean solid loading per slice
- Represent unresolved pathways with a single diffusivity D

Fit D so the reduced model reproduces the network's mean uptake history.

The one-dimensional cells do not come from the pore-network extraction. We create a separate grid by dividing bed height into one hundred equal slices spanning the tube cross-section. Each slice contains homogenized gas and adsorbent phases. The detailed pore paths disappear from this representation. Their impact on transient uptake is represented by one fitted diffusivity. The reduced system has two hundred concentration and loading states, while the PNM has a state for every pore and bead. We compare their mean uptake curves rather than attempting a one-to-one mapping of their internal fields.

The homogeneous 1D uptake model

Gas balance in each axial cell

$$\varepsilon_b \frac{\partial c}{\partial t} = D \frac{\partial^2 c}{\partial z^2} - (1 - \varepsilon_b) \rho_p \frac{\partial q}{\partial t}.$$

Adsorbent loading in each axial cell

$$\frac{\partial q}{\partial t} = k_{\text{LDF}} \left[\frac{q_s b R T c}{1 + b R T c} - q \right].$$

- Divide the bed height into 100 equal slices, each spanning the tube cross-section.
- Solve gas concentration and mean solid loading together in each slice.
- Trial D changes c , which changes adsorption.

The 1D grid is separate from the pore network.

Average the calculated loading

$$\bar{q}_{1D}(t; D) = \frac{1}{H} \int_0^H q(z, t; D) dz.$$

For 100 equal axial cells:

$$\bar{q}_{1D}(t; D) = \frac{1}{100} \sum_{\ell=1}^{100} q_{\ell}(t; D).$$

- Match bed height, porosity and inlet condition to the network.
- Keep isotherm, kinetics and initial state fixed. Only D changes.

2026-09-12

Reduction and diffusivity prediction

The homogeneous 1D uptake model

This slide describes the homogeneous one-dimensional model, not the individual pore network. Each axial cell contains gas and a homogenized amount of adsorbent. We solve one gas concentration and one solid loading per cell, giving 200 coupled state equations for 100 cells. The loading equation explicitly substitutes the Langmuir equilibrium relation and p equals cRT . The diffusivity appears in the gas balance. Changing it changes the gas concentration history and therefore the adsorption history. The right-hand side is a spatial average of the calculated loadings, not an additional independent evolution equation. For equal cells and uniform adsorbent density, the arithmetic average is the bed mean loading.

The homogeneous 1D uptake model

Gas balance in each axial cell

$$\varepsilon_b \frac{\partial c}{\partial t} = D \frac{\partial^2 c}{\partial z^2} - (1 - \varepsilon_b) \rho_p \frac{\partial q}{\partial t}.$$

Adsorbent loading in each axial cell

$$\frac{\partial q}{\partial t} = k_{\text{LDF}} \left[\frac{q_s b R T c}{1 + b R T c} - q \right].$$

For 100 equal axial cells:

$$\bar{q}_{1D}(t; D) = \frac{1}{100} \sum_{\ell=1}^{100} q_{\ell}(t; D).$$

- Divide the bed height into 100 equal slices, each spanning the tube cross-section.
- Solve gas concentration and mean solid loading together in each slice.
- Trial D changes c , which changes adsorption.
- Match bed height, porosity and inlet condition to the network.
- Keep isotherm, kinetics and initial state fixed. Only D changes.

The 1D grid is separate from the pore network.

Fit the effective diffusivity to the complete uptake curve

Per-bed optimization

1. Select a trial D within prescribed bounds.
2. Solve the coupled 1D equations from the same initial state.
3. Compare the full uptake curve with the PNM reference.
4. Use previous errors to select the next D and repeat (optimization).

Curve error and fitted coefficient

$$e_m(D) = \bar{q}_{1D}(t_m; D) - \bar{q}_{PNM}(t_m).$$

$$NRMSE(D) = \frac{\sqrt{\sum_m w_m e_m(D)^2}}{\bar{q}_{PNM}(t_f)},$$

$$\sum_m w_m = 1.$$

$$D_{eff} = \arg \min_D NRMSE(D).$$

The time weights account for unequal spacing of saved outputs.

A fitted D_{eff} is a reduction target for each bed, not yet a geometry-based prediction.

For each bed, we repeatedly solve the homogeneous model with different diffusivities. NRMSE means normalized root-mean-square error. Here we weight squared errors by elapsed time using trapezoidal weights, take the square root and divide by the final reference loading. This prevents densely saved times from dominating the fit. The bounded scalar optimizer operates in $\log_{10} D$; it starts deterministically and previous errors guide later trials. Only D changes. This produces one reduction target per packing. The next question is whether a cheaper calculation from geometry can predict that target.

Can inlet conductance predict the fitted diffusivity?

Calculate steady inlet access

- Use the pore network with adsorption disabled.
- Fix concentrations at the inlet and an interior plane.
- Solve steady diffusion and sum the molar flux $\dot{n}$.

$$G_{\text{access}} = \frac{\dot{n}}{\Delta c} \quad [\text{m}^3/\text{s}].$$

Conductance measures how much gas the inlet-connected pathways can supply per concentration difference.

One steady conductance calculation supplies the input for transient prediction.

Calibrate the transport coefficient

- Pair each calculated G_{access} with its fitted D_{eff} .
- Fit $D_{\text{eff}} = f(G_{\text{access}})$ on training beds.
- Predict the omitted bed's diffusivity.
- Run its 1D model and compare the complete uptake curve.

└ Reduction and diffusivity prediction

└ Can inlet conductance predict the fitted diffusivity?

Can inlet conductance predict the fitted diffusivity?

Calculate steady inlet access

- Use the pore network with adsorption disabled.
- Fix concentrations at the inlet and an interior plane.
- Solve steady diffusion and sum the molar flux $\dot{n}$.

$$G_{\text{access}} = \frac{\dot{n}}{\Delta c} \quad [\text{m}^3/\text{s}].$$

Conductance measures how much gas the inlet-connected pathways can supply per concentration difference.

One steady conductance calculation supplies the input for transient prediction.

Calibrate the transport coefficient

- Pair each calculated G_{access} with its fitted D_{eff} .
- Fit $D_{\text{eff}} = f(G_{\text{access}})$ on training beds.
- Predict the omitted bed's diffusivity.
- Run its 1D model and compare the complete uptake curve.

G access is calculated from a steady diffusion solve on the network, not predicted. Total molar flux divided by the imposed concentration difference has units of cubic metres per second. We use this conductance directly as the descriptor. A relationship learned from other beds maps it to the effective diffusivity in the transient 1D model. Conductance can also correlate with final uptake, but here its target is the coefficient fitted to the complete uptake history. For validation we first predict the omitted bed diffusivity and then run its 1D model without fitting to that bed history. The optional normalization appendix presents the alternative descriptor.


- Reduction and diffusivity prediction
- Why use an interior plane?

Why use an interior plane?



- A plane at $5d_p$ includes**
- the inlet openings
 - second-layer and deeper bottlenecks
 - parallel paths connecting inlet to the plane
- The bed top answers a different question**
- measure conductance through the entire bed
 - include upper bottlenecks that may matter more at later times

Why use an interior plane?



A plane at $5d_p$ includes

- the inlet openings
- second-layer and deeper bottlenecks
- parallel paths connecting inlet to the plane

The bed top answers a different question

- measure conductance through the entire bed
- include upper bottlenecks that may matter more at later times


The plane provides a tunable spatial scale for the accessibility diagnostic. A near-inlet plane includes the entrance openings and some deeper constrictions and parallel paths. Using the bed top instead measures conductance through the complete network. The focus on a near-inlet region is motivated by entry-driven uptake, but it does not establish that upper pores are irrelevant at early times. In the current implementation, depth is measured above the lowest retained pore. Five bead diameters is the selected diagnostic depth. The results also compare three, seven and ten bead diameters to examine sensitivity.

└ Reduction and diffusivity prediction


└ Development and later use are two different workflows

Development and later use are two different workflows

Develop closure




Predict for a new bed



For a new bed, the transient pore-network adsorption model is no longer required.

During development each bed contributes a steady calculated inlet conductance and an effective diffusivity fitted to its complete transient PNM uptake curve. We calibrate the relation between them. For a new bed we calculate its steady conductance, predict its diffusivity, and run the inexpensive 1D transient model. The new bed still requires geometric quantities such as height and porosity for that model, but no transient PNM uptake solve is required for prediction. Held-out evaluation uses the PNM curve only afterwards to assess error.




Section 4 of 5

Results and robustness

Test the physical hypothesis and the model reduction

Test the physical hypothesis and the model reduction.

The twenty beds have nearly equal bulk porosity



Bulk-porosity CV: 0.55%.

Controlled ensemble

- same bead count and diameter
- same tube geometry
- only the insertion seed changes

CV: coefficient of variation

$$CV_{\epsilon} = s_{\epsilon} / \bar{\epsilon}$$

s_{ϵ} : standard deviation


$\bar{\epsilon}$: mean porosity

2026-09-12

Results and robustness


The twenty beds have nearly equal bulk porosity

The twenty beds have nearly equal bulk porosity



The mean bulk porosity is approximately 0.44665. Divide the sample standard deviation by this mean to obtain the coefficient of variation, 0.55 percent. The horizontal axis orders beds by their insertion seed. The small spread establishes the controlled comparison.

Similar porosity still produces different uptake curves



At 5000 s: 0.206–0.315 mol kg⁻¹, with an uptake CV of 11.0%.

Network reference


- all twenty uptake curves
- highlighted beds span low, median and high final uptake
- identical adsorption parameters

2026-09-12

Results and robustness


Similar porosity still produces different uptake curves

Similar porosity still produces different uptake curves



The pale curves show the full ensemble. Orange, blue and green identify low, median and high final uptake. These representative beds are selected by final uptake, not by accessibility. The uptake varies much more strongly than porosity.

Bulk porosity alone does not explain final uptake



Porosity baseline

- each point represents one bed
- weak rank association
- wide uncertainty around the fitted relationship


Void volume alone does not describe access to the adsorbent.

2026-09-12

Results and robustness


Bulk porosity alone does not explain final uptake

Bulk porosity alone does not explain final uptake



The plot compares bulk porosity with mean loading at 5000 seconds. Spearman correlation is approximately 0.102. The shaded band is pointwise bootstrap uncertainty in the fitted relationship, not a prediction interval for a new bed.

Accessibility explains differences in final uptake



Inlet-connected pathways explain variability that porosity misses.

Steady network descriptor

- G_{access} measures inlet-to-interior conductance
- greater conductance accompanies greater uptake


Spearman $\rho_s = 0.968$.

2026-09-12

Results and robustness


Accessibility explains differences in final uptake

Accessibility explains differences in final uptake



We calculate accessibility using a steady diffusion problem without adsorption. The Spearman coefficient is 0.968. More accessible beds generally have greater uptake. This is a within-ensemble association. The shaded band represents uncertainty in the fitted relationship.

Greater accessibility reduces particle underutilization



Higher G_{access} accompanies fewer underutilized particles.

Underutilized fraction

$$f_{\text{under}} = \frac{1}{N_p} \sum_{k=1}^{N_p} I_k.$$

$I_k = 1$ when

$$q_k(t_f) < 0.5q^*(p_{\text{inlet}}).$$


$t_f = 5000 \text{ s.}$

The indicator equals one for a particle below half the inlet-equilibrium loading.

2026-09-12


Results and robustness

Greater accessibility reduces particle underutilization



For each bead we compare its final loading with half the equilibrium loading at the inlet pressure. The indicator is one below this threshold and zero otherwise. Averaging over particles gives the plotted fraction. This is a threshold-based utilization measure, not the fraction of completely inaccessible particles.

One fitted D_{eff} reproduces each uptake curve



Across twenty beds, the mean fitted-curve NRMSE is 1.58%.

Per-bed optimization

1. solve the 1D model for a trial D
2. evaluate time-weighted NRMSE against the network curve
3. use previous errors to select the next D and repeat (optimization)


Solid: network.
Dashed: fitted 1D model.

2026-09-12

Results and robustness

One fitted D_{eff} reproduces each uptake curve

One fitted D_{eff} reproduces each uptake curve




Per-bed optimization

1. solve the 1D model for a trial D
2. evaluate time-weighted NRMSE against the network curve
3. use previous errors to select the next D and repeat (optimization)

Solid: network.
Dashed: fitted 1D model.

For each bed, the transient pore-network calculation provides a reference uptake curve. We repeatedly solve the 1D model with trial diffusivities and minimize the time-weighted normalized root-mean-square error. The bounded optimizer searches in $\log_{10} D$. Its initial trial follows from the prescribed bounds, without random sampling. Previous objective values guide subsequent trials until the search meets its tolerance. Each trial solves both coupled equations from the same initial state. The optimized diffusivity is D_{eff} . This figure shows per-bed fits, not predictions for unseen beds. The mean error is 1.58 percent.

Residuals show the remaining reduction error



Time-weighted NRMSE across all beds: 0.84–2.46%.

Curve residual

$$e(t) = \bar{q}_{1D}(t) - \bar{q}_{PNM}(t).$$


- positive: excess predicted uptake
- negative: insufficient predicted uptake

2026-09-12

Results and robustness


Residuals show the remaining reduction error

Residuals show the remaining reduction error



The same three representative beds appear here. The residual is the 1D mean loading minus the network mean loading. Its sign changes over time, showing that one coefficient cannot reproduce every detail. NRMSE summarizes the full residual curve using trapezoidal time weights and normalizes by final network loading.

Three relationships use the same inlet conductance



Predict D_{eff} from G_{access}

Proportional: $D_{\text{eff}} = \alpha_G G_{\text{access}}$

Linear: $D_{\text{eff}} = a_G G_{\text{access}} + b_G$

Power law: $D_{\text{eff}} = C_G G_{\text{access}}^n$

All coefficients are calibrated.
Their units make D_{eff} a diffusivity.


The power-law exponent is $n \approx 2.52$ for these twenty beds.

The predictor stays fixed; the comparison tests the relationship's form.

2026-09-12


Results and robustness

Three relationships use the same inlet conductance



This figure uses raw inlet conductance for all three relationships. The proportional form has one coefficient. The linear and power-law forms have two. The power law is fitted by least squares in logarithmic coordinates, while linear and proportional fits use diffusivity coordinates; all held-out scores are then evaluated in physical diffusivity units. The fitted exponent is about 2.52. Coefficients carry the units required by the equations, so their numerical values depend on the unit convention. This full-ensemble plot illustrates the relationships, not their held-out accuracy. The power law is positive for positive conductance and approaches zero as conductance goes to zero. The linear fit has a negative intercept, so extrapolation outside the observed conductance range can become nonphysical. Neither form is established as universal.

Conductance predicts diffusivity for an omitted bed



Leave-one-bed-out comparison


Predictor / form	R^2
G_{access} / power law	0.906
G_{access} / linear	0.898
G_{access} / proportional	0.558
Porosity / linear	-0.234

Fit every coefficient on nineteen beds; predict the omitted bed.
Repeat for all twenty beds.

Power-law and linear predictions are close; both outperform porosity in this ensemble.


2026-09-12

- Results and robustness
 - Conductance predicts diffusivity for an omitted bed



All conductance-based forms now use the same G access predictor. We omit one bed, fit all coefficients to the other nineteen and predict the omitted fitted diffusivity. Power-law R squared is 0.906 and linear R squared is 0.898. This small difference does not establish decisive superiority. Proportionality gives 0.558 and the porosity baseline gives minus 0.234. The diffusivity target itself was obtained from the omitted PNM history; it is withheld from calibration of the relationship. Model form and depth were considered within this ensemble, so these scores are internal leave-one-bed-out validation, not independent validation of model selection. We next test the actual uptake histories generated from those predicted coefficients.

Predicted diffusivity reproduces the uptake history



Mean curve NRMSE, twenty beds

Relationship	Error
G_{access} / power law	2.61%
G_{access} / linear	2.64%
G_{access} / proportional	4.80%
Porosity / linear	7.58%

Use each held-out D_{eff} in the 1D solver, with no refitting.


Time-weighted RMSE divided by the bed’s final PNM uptake.

Conductance predicts a transport coefficient and therefore the full history, not only final uptake.

2026-09-12


Results and robustness

Predicted diffusivity reproduces the uptake history



We now evaluate the end-to-end prediction. For each omitted bed, its predicted diffusivity is supplied to the same 100-cell 1D solver with the same material properties, inlet boundary, initial state, height and porosity as the original fit. We do not refit that coefficient to its uptake curve. All four relationships are propagated for all twenty beds, giving eighty simulations. The plotted solid curves are PNM references for low, median and high final uptake; dashed curves use held-out power-law predictions. The table reports the mean error across all twenty beds for every model. Curve NRMSE weights time by trapezoidal integration and divides by final reference uptake. Power law and linear give nearly identical mean errors, 2.61 and 2.64 percent. These include both coefficient prediction error and 1D reduction error. The earlier 1.58 percent measures the per-bed fitted reduction error and is a different quantity. G access can also be regressed directly against final uptake, but this route instead predicts D eff to obtain the complete transient response.

Conductance correlation across diagnostic depths



Depth sensitivity

- calculate G_{access} to planes at 3, 5, 7, $10d_p$
- Pearson: linear association
- Spearman: rank association


Depth is measured above the network's lowest pore.

The conductance–diffusivity association persists across the tested interior planes.

2026-09-12


Results and robustness

Conductance correlation across diagnostic depths



This is the direct conductance descriptor, recalculated from the inlet to each of four interior planes. We compare each set of conductances with the same per-bed fitted effective diffusivities. Pearson measures linear association and Spearman measures monotonic rank association. The correlations remain high across these depths. This is an association check; the next slide refits and validates the predictive relationships at each depth.

Conductance-model performance across diagnostic depths



Four curves at every depth

$$D_{\text{eff}} = C_G G_{\text{access}}^n$$

$$D_{\text{eff}} = a_G G_{\text{access}} + b_G$$

$$D_{\text{eff}} = \alpha_G G_{\text{access}}$$

$$D_{\text{eff}} = a_0 + a_1 \varepsilon_b$$

Refit on nineteen beds at each depth; predict the omitted bed.


The porosity baseline is independent of plane depth.

Power-law and linear conductance models remain predictive across the tested depths.

2026-09-12

Results and robustness

Conductance-model performance across diagnostic depths



The figure contains four curves: power-law, linear-with-intercept and proportional relations using G_{access} , plus the porosity baseline. Every conductance relation is refitted within every leave-one-bed-out fold at each depth. The porosity curve stays constant because the predictor does not depend on the chosen plane. Five particle diameters gives the best power-law and linear score among these four planes for this ensemble. This establishes sensitivity within our existing data; it does not establish an optimal plane for other geometries or independently validate the model and depth choice. The normalized alternatives are kept in the separate optional appendix.

Section 5 of 5

Conclusions and next steps

What is established and what remains to validate

What is established and what remains to validate.

What the present evidence establishes

Supported

- porosity alone misses important packing variability
- inlet-connected topology controls transient utilization
- one D_{eff} reproduces each detailed curve accurately
- G_{access} provides an interpretable predictive closure

Not yet claimed

- industrial breakthrough performance
- effects of flow, heat or water competition
- transfer to other bead sizes or D_{tube}/d_p
- a universal conductance–diffusivity relation

The contribution is a structure-to-model closure for a controlled diffusion-driven adsorption problem.

2026-09-12

└─ Conclusions and next steps

└─ What the present evidence establishes

What the present evidence establishes


Supported	Not yet claimed
• porosity alone misses important packing variability • inlet-connected topology controls transient utilization • one D_{eff} reproduces each detailed curve accurately • G_{access} provides an interpretable predictive closure	• industrial breakthrough performance • effects of flow, heat or water competition • transfer to other bead sizes or D_{tube}/d_p • a universal conductance–diffusivity relation

The contribution is a structure-to-model closure for a controlled diffusion-driven adsorption problem.

The evidence supports a steady inlet-conductance descriptor that captures transport differences missed by bulk porosity. A single fitted effective diffusivity reproduces each transient PNM curve with a mean error of 1.58 percent. Using conductance to predict that diffusivity gives a mean held-out curve error of 2.61 percent with the power law, essentially comparable to 2.64 percent for a linear model. The result is specific to this controlled diffusion-adsorption ensemble. We do not claim industrial breakthrough performance, robustness to heat or water competition, or transfer across bead sizes, tube geometries, sorbents and operating conditions. A more general study would be needed to establish those claims.

Pore accessibility retains what bulk porosity loses

Equal void volume does not imply equal transport. How that void space connects to the inlet matters.



2026-09-12

Conclusions and next steps

Pore accessibility retains what bulk porosity loses

Pore accessibility retains what bulk porosity loses

Equal void volume does not imply equal transport. How that void space connects to the inlet matters.



The central result is that inlet-connected gas conductance supplies information that porosity misses. Starting from a packing we calculate the network and its steady conductance, predict the effective diffusivity from a calibrated relationship, and run the homogeneous 1D transient model. The twenty-bed ensemble provides a fitted reduction error of 1.58 percent and a power-law held-out uptake-history error of 2.61 percent. These describe different steps in the workflow. The power-law diffusivity R squared is 0.906 and is close to the linear model score of 0.898. The aim is a simpler interpretable predictor for the controlled problem, not a claim of universal model superiority.